DATE
Thursday, Feb. 5, 2026 at 5 p.m. ET

CALL PARTICIPANTS
- Chief Executive Officer — Andrew Jassy
- Chief Financial Officer — Brian T. Olsavsky
- Director, Investor Relations — Dave Fildes

TAKEAWAYS
- Revenue -- $213.4 billion, representing a 12% year-over-year increase excluding foreign exchange impact.
- Operating Income -- $25 billion, includes $2.4 billion in special charges for tax disputes, severance costs, and asset impairments.
- Free Cash Flow (Trailing Twelve Months) -- $11.2 billion, directly stated by management.
- AWS Revenue -- $35.6 billion, up 24% year over year; highest growth rate in thirteen quarters.
- AWS Annualized Run Rate -- $142 billion; AWS added $2.6 billion sequentially and nearly $7 billion year over year.
- AWS Backlog -- $244 billion, up 40% year over year and 22% sequentially.
- AWS Operating Income -- $12.5 billion with a 35% margin, up 40 basis points year over year (cited by Olsavsky).
- Custom Chips Revenue Run Rate -- Over $10 billion annually, with Graviton and Trainium units growing at triple-digit percentages and more than 50% year over year for Graviton.
- Bedrock AI Platform -- Multibillion-dollar annualized run rate, with customer spend up 60% quarter over quarter.
- Prime Delivery Speeds -- U.S. Prime members received over 8 billion items same or next day, growing more than 30% year over year.
- Third-Party Seller Unit Mix -- 61% of worldwide paid units, with continued investment in seller-facing AI tools.
- Advertising Revenue -- $21.3 billion for the quarter, up 22% year over year; $12 billion in incremental revenue for the year.
- Amazon LEO Satellite Network -- 180 satellites launched to date; more than 20 launches planned for 2026 and over 30 in 2027, with commercial rollout expected later in 2026 and several enterprise customers already signed.
- Q1 2026 Guidance -- Net sales expected between $173.5 billion and $178.5 billion, with a favorable 180 basis-point foreign exchange impact.
- Q1 2026 Operating Income Guidance -- Operating income projected between $16.5 billion and $21.5 billion.

Need a quote from a Motley Fool analyst? Email pr@fool.com

RISKS
- Olsavsky highlighted a projected $1 billion year-over-year cost increase in North America during Q1 2026 due to higher Amazon LEO satellite expenditures, noting, "we are spending more on launching satellites each year."
- The call cited a $2.4 billion aggregate impact from special charges—tax disputes, severance costs, and asset impairments—reducing reported operating income.

SUMMARY
Amazon **(AMZN2.48%)** reported 12% revenue growth and accelerated AWS expansion, supported by new customer adoption and an expanded services backlog. Management called out AI and custom silicon as rapidly scaling businesses, substantiated by substantial sequential and annual gains. The company increased satellite deployment and fulfillment speed, aligning with rising investments in cloud infrastructure and global retail logistics.

- Amazon's executive team disclosed that Bedrock's customer spend rose 60% sequentially and that Rufus, the AI agent, enabled over 300 million customer uses, driving a 60% lift in purchase likelihood among users.
- North America segment revenue was $127.1 billion (+10%) and International segment revenue reached $50.7 billion (+11% excluding FX), with International segment margins expanding when excluding special charges.
- Prime Video's global ad-supported audience grew to 315 million from 200 million at the start of 2024, evidencing traction in streaming ad monetization.
- Management stated that in 2025, AWS added more data center capacity than any other company in the world.
- Amazon Hall surpassed one million sub-$10 items in 25 countries, and everyday essentials represented one-third of U.S. store unit sales, underlining expansion in value-driven categories.

INDUSTRY GLOSSARY
- Graviton and Trainium: Amazon-designed custom silicon chips used for general cloud compute (Graviton) and AI-specific workloads (Trainium), cited as driving price and performance gains in AWS.

- Bedrock: Amazon's managed service for deploying foundation AI models, inference, and fine-tuning, offered to AWS customers on demand.
- Prime Video ad-supported audience: The global population of viewers reached via advertising-supported Prime Video content.
- Amazon LEO: Amazon's Low Earth Orbit satellite constellation for global broadband and enterprise connectivity.
- Agentic AI/Agent Corp/Frontier Agents: Amazon's terminology for autonomous or semi-autonomous AI applications that perform tasks for users or enterprises, within AWS or retail services.

Full Conference Call Transcript

Dave Fildes: Hello, and welcome to our Q4 2025 financial results conference call. Joining us today to answer your questions is Andrew Jassy, our CEO, and Brian T. Olsavsky, our CFO. As you listen to today's conference call, we encourage you to have our press release in front of you, which includes our financial results as well as metrics and commentary on the quarter. Please note, unless otherwise stated, all comparisons in this call will be against our results for the comparable period of 2024. Our comments and responses to your questions reflect management's views as of today, February 5, 2026, only and will include forward-looking statements. Actual results may differ materially.

Additional information about factors that could potentially impact our financial results is included in today's press release and our filings with the SEC, including our most recent annual report on Form 10-K and subsequent filings. During this call, we may discuss certain non-GAAP financial measures. In our press release, slides accompanying this webcast, and our filings with the SEC, each of which is posted on our IR website, you will find additional disclosures regarding these non-GAAP measures, including reconciliations of these measures with comparable GAAP measures. Our guidance incorporates the order trends that we have seen today and what we believe today to be appropriate assumptions.

Our results are inherently unpredictable and may be materially affected by many factors, including fluctuations in foreign exchange rates and energy prices, changes in global economic and geopolitical conditions, tariff and trade policies, resource and supply volatility, including for memory chips, and customer demand and spending, including the impact of recessionary fears, inflation, interest rates, regional labor market constraints, world events, the rate of growth of the Internet, online commerce, cloud services, and new and emerging technologies, and the various factors detailed in our filings with the SEC. Our guidance assumes, among other things, that we do not conclude any additional business acquisitions, restructurings, or legal settlements.

It's not possible to accurately predict demand for our goods and services, and therefore, our actual results could differ materially from our guidance. And now I'll turn the call over to Andrew Jassy.

Andrew Jassy: Thanks, Dave. We are reporting $213.4 billion in revenue, up 12% year over year excluding the impact from foreign exchange rates. Operating income was $25 billion, and trailing twelve-month free cash flow was $11.2 billion. We are seeing strong growth, and with the incremental opportunities available to us in areas like AI, chips, low earth orbit satellites, quick commerce, and serving more consumers' everyday essentials needs, we have a chance to build an even more meaningful business in Amazon.com, Inc. in the coming years. With strong return on invested capital, and we are investing to do so. We are already seeing strong demand in these areas even in these early innings. I'll start with AWS.

AWS growth continued to accelerate to 24%, the fastest we've seen in thirteen quarters. Up $2.6 billion quarter over quarter and nearly $7 billion year over year. AWS is now a $142 billion annualized run rate business. And our chips business, inclusive of Graviton and Trainium, is now over $10 billion in annual revenue run rate growing triple-digit percentages year over year. As a reminder, it's very different having 24% year-over-year growth on a $142 billion annualized run rate than to have a higher percentage growth on a meaningfully smaller base which is the case with our competitors. We continue to add more incremental revenue and capacity than others, and extend our leadership position.

We are continuing to see strong growth in core non-AI workloads as enterprises return to focusing on moving infrastructure from on-premises to the cloud, along with AWS having the broadest functionality, strongest security and operations performance, and most vibrant partner ecosystem. AWS continues to earn most of the big enterprise and government transitions to cloud. Since our last call, we announced new agreements with OpenAI, Visa, MBA, BlackRock, Perplexity, Lyft, United Airlines, DoorDash, Salesforce, US Air Force, Adobe, Thomson Reuters, AT&T, S&P Global, National Bank of Canada, the London Stock Exchange, Choice Hotels, Accenture, Indeed, HSBC, CrowdStrike, and many more. More of the top 500 US startups use AWS as their primary cloud provider than the next two providers combined.

We are adding significant EC2 core computing capacity each day, the majority of that new compute is using our custom CPU silicon Graviton. Graviton is up to 40% more price per than leading x86 processors and is used expansively by over 90% of

AWS's top thousand customers. Graviton itself is a multibillion-dollar annualized run rate business growing more than 50% year over year. We consistently see customers wanting to run their AI workloads where the rest of their applications and data are. We are also seeing that as customers run large AI workloads on AWS, they are adding to their core AWS footprint as well.

But the biggest reason that AWS continues to gain AI share is our uniquely broad top-to-bottom AI stack functionality. In AI, we are doing what we've always done in AWS, solving customer challenges. Let me give you some examples. The first challenge is having a strong foundation model to generate inferences or predictions. Customers are realizing as they get further into AI that they need choice. As different models are better on different dimensions. In fact, most sophisticated AI applications leverage multiple models. Whether customers want frontier models like Anthropix Cloud, or open models like Miesztrall or Lama, Frontier Intelligence will lower cost and latency like Amazon Nova. Or video and audio models like twelve Labs or NovaSonic.

Amazon Bedrock makes it easy to use these models to run inference scalably, and performantly. Bedrock is now a multibillion-dollar annualized run rate business, and customer spend grew 60% quarter over quarter. The second challenge is how to hone the model for your application. Customers sometimes think if they have a good model, they will have a good AI application. It's not really true. It takes a lot of work to post-train and fine-tune a model for your application. Our SageMaker AI service along with fine-tuning tools in Bedrock make this much easier for customers. A third challenge is how to have a custom version of a foundation that best leverages the company's secret sauce, their own data.

To date, companies have tried to shape models with their own data late in the process, usually with fine-tuning or post-training. There's debate in the industry about this, but we believe that enterprises will want models trained on their own data at an early stage, at pre-training if possible, so their models have the best possible foundation for what matters most to each enterprise on which to learn and evolve. It's a little like teaching a child a foreign language early in their life. That becomes part of their learning foundation moving forward and it makes it easier to pick up other languages later in their life. To solve for this need, we just launched NovaForge.

Which gives customers early checkpoints on our Amazon Nova models allows them to securely mix their own proprietary data with the model's data in the pre-training stage, and enables their own uniquely customized versions of Novo. What we call novellas, trained with their data early in the process. This will be very useful for companies as they build their own agents on top of the model. There is nothing else out there like this today and a potential game changer for companies. Another challenge is cost. I've said this many times, but if we want AI to be used as expansively as companies want, we have to make the cost of inference lower.

A significant impediment today is the cost of AI chips. Customers are starving for better price performance and typically and understandably, the dominant early leaders aren't in a hurry to make that happen. They have other priorities. It's why we built our own custom silicon in training. And it's really taken off. We've landed over 1.4 million Trainium two chips our fastest ramping chip launch ever. Trainium two is 30 to 40% more price performance than comparable GPUs. It is a multibillion-dollar annualized revenue run rate business with a 100,000 plus companies using it is trading as the majority underpinning of Bedrock usage today. We recently launched Trainium three, which is up to 40% more price performance than Trainium two.

We are seeing very strong demand for Trainium three and expect nearly all of our Trainium three supply of chips to be committed by mid-2026. And though we are still building Trainium four, we are seeing very strong interest already. Looking ahead, the primary way companies will get value from AI is with agents. Some their own, some from others, and there are several customer challenges that we are well positioned to solve. It's harder to build agents than it should be. For that, we've built strands of service enabling agents to be created from any model.

Once agents are built, enterprises are apprehensive about deploying to production because these agents need to securely and scalably connect to compute, data tools, memory, identity, policy governance, performance monitoring, and other elements. This is a new and hard problem where a solution has not existed until we launched Bedrock Agent Corp. Customers are quite excited about Agent Core, and it's unlocking deployments. Customers also want to leverage others' useful agents, and we've built several. Including Curo for coding, Amazon Quick for knowledge workers to leverage their own data and analytics, AWS Transform for software migration, and Amazon Connect for call center operations. We continue adding new capabilities, and usage continues to grow quickly.

For example, the number of developers using Curo grew more than 150% quarter over quarter. In addition to agents that customers direct, customers are also becoming excited about agents that require less human interaction. They can be fully autonomous, run persistently for hours or days, scale out quickly, and remember context. At this past AWS re:Invent, we launched Frontier Agents to do that. Kiro autonomous agents for coding tasks, AWS DevOps agents for detecting and resolving

operational issues, and AWS security agents for proactively securing applications throughout the development life cycle, and they're already making a big difference for customers.

We expect to invest about $200 billion in capital expenditures across Amazon.com, Inc., but predominantly in AWS, because we have very high demand. Customers really want AWS for core and AI workloads. And we are monetizing capacity as fast as we can install it. We have deep experience understanding demand signals in the AWS business and then turning that capacity into strong return on invested capital. We are confident this will be the case here as well. I'll now turn to stores.

We continue to expand selection, including more 400 new beauty brands in The US in 2025, like Bobbi Brown Cosmetics, Charlotte Tilbury, and Laura Mercier, and new fashion brands like Away Luggage, Converse, Diesel, Michael Kors, Nike, and The North Face. Our ultra-low priced offering, Amazon Hall, grew selection to over a million items under $10 in expanded to serve customers in more than 25 countries and regions. We continue to see strong customer response to everyday essentials and grocery. In 2025, everyday essentials grew nearly twice as fast as all other categories in The US, representing one out of three units sold in our store. And we've become a go-to grocery destination for over 150 million Americans.

Mostly through online shopping and Whole Foods. With over $150 billion in gross sales, Amazon.com, Inc. is clearly a large grocer at this point. Customers in thousands of US cities and towns can now get perishables delivered same day alongside millions of other items. And customers who use that service shop more than twice as often as customers who don't. We plan to expand in many more communities in 2026, and we also plan to open more than 100 new Whole Foods Market stores over the next few years as we work to make grocery shopping easier, and more affordable for customers. We remain committed to staying sharp on price. And continue to meet or beat other retailers' prices.

A recent study from Profitero showed that Amazon.com, Inc. is America's lowest priced retailer for the ninth straight year. 14% lower on average than other major online retailers. We again achieved our fastest ever delivery speeds for Prime members around the world in 2025. In The US, we delivered nearly 70% more items same day than the year before. We also continue increasing speed for rural customers with nearly two times more average monthly customers in rural areas receiving same day delivery year over year. Same day is our fastest growing delivery offering and nearly 100 million customers used it last year in The US.

And the team is continuing to innovate, We've launched Amazon now in India, Mexico, and The UAE. Ultrafast delivery on thousands of items in about thirty minutes or less. And we are testing it in several communities in The US and UK. It's early, but customers are loving it. For example, in India, customer response exceeded our most optimistic expectations, and we are seeing Prime members triple their shopping frequency once they start using it. Expanding our same day delivery coverage also leads to meaningfully later cutoff times for orders. Which is a big deal for customers. For example, on Christmas Eve, customers in about 4,000 US cities could order items up until midday and get them that same day.

Another example is our recently launched feature add to delivery. Which enables Prime members in The US to add items to their upcoming Amazon.com, Inc. deliveries with just one tap without going through checkout again or paying additional shipping fees. Just six months after launch, add to delivery already makes up about 10% of all Prime volume fulfilled through the Amazon.com, Inc. network each week. While this seems simple on the surface, this feature is supported by a lot of invention where we need to figure out in real time with incredibly low latency what items among Amazon.com, Inc.'s hundreds of millions of products are available to add to a customer's upcoming deliveries?

Surface them, find a way to include in their packages, and deliver within the same customer promise. The stores team is also continuing to innovate and deliver for customers with AI. Our AgenTik AI shopping assistant, Rufus, is rapidly expanded. Rufus can research products, track prices, and auto buy. Purchasing a product in our store when it reaches your set price. It can also now shop tens of millions of items in other online stores and make purchases for customers using our agentic buy for me feature. Last year, more than 300 million customers used Rufus. In addition, customers used Lens our AI powered visual search tool to find products with a phone's camera, a screenshot, or a barcode.

And they did it 45% more year over year. Moving on to Amazon Ads. We are pleased with the continued strong growth across our full funnel offerings generating $21.3 billion of revenue in the quarter and growing 22% year over year. Sponsored products advertising in our store continues to be our largest ads offering, and the combination of trillions of shopping, browsing, and streaming signals with advanced AI and machine learning led us deliver highly relevant useful ads for customers. We saw continued growth in Prime Video ads. Which is now available in 16 countries and is contributing meaningfully to our revenue growth.

Prime Video has an average ad supported audience of 315 million viewers globally, up from 200 million in early 2024. Our ads team is also innovating with AI. We recently announced our ads agent, which lets brands use AI to create non optimize campaigns at scale. Implement effective campaign targeting, and quickly create actionable insights. And our creative agent lets advertisers research brainstorm, and generate full funnel ad campaigns from concept to completion using conversational guidance in Amazon.com, Inc.'s retail data transforming what was a weeks long process into just hours. We are also continuing to invent and see momentum several other areas, and I'll mention just a few. Starting with live sports on Prime.

The fourth season of Thursday night football broke more records. It was our most watched season ever, averaging more than 15 million viewers, a 16% year over year increase, and a third consecutive year of double digit growth. And the packers versus bears wild card game was the most streamed NFL game in history with 31.6 million viewers clearing the prior mark by more than 4 million. We just made Alexa Plus available to all customers in The US, free for prime members and $19.99 a month for non prime members.

Alexa Plus continues to get even better and more capable and we've added new ways to interact with Alexa, including a new chat experience at alexa.com, a redesigned mobile app, and new integrations with third party devices like Samsung TVs and BMW cars. We've also added new features like the ability to answer a ring doorbell on a customer's behalf, and more ways to shop or manage a home. And finally, the team is making rapid progress on Amazon LEO. Which will bring connectivity to consumers, enterprises, and governments in places where they don't have broadband connectivity. Our enterprise grade customer terminal, LEO Ultra, is the fastest satellite Internet antenna ever built.

Delivering simultaneous download speeds of up to one gigabit per second upload speeds of up to 400 megabits per second. LEO will offer enterprise grade performance and advanced encryption with secure private networking that bypasses public Internet. Connecting directly to AWS. We've launched 180 satellites have more than 20 launches planned in 2026, more than 30 in 2027, and expect to launch commercially in 2026. We have dozens of commercial agreements already signed, including with AT&T, DIRECTV Latin America, JetBlue, and Australia's national broadband network. I have many more on the way. It's been an action-packed year of innovation and progress and we've hit the ground running in 2026.

With that, I'll turn it over to Brian for a financial update.

Brian T. Olsavsky: Thanks, Andy. Starting with our top line financial results. Worldwide revenue was $213.4 billion, a 12% increase year over year excluding the 150 basis points favorable impact of foreign exchange. In Q4, we reported worldwide operating income of $25 billion. This operating income includes three special charges, which reduced operating income by $2.4 billion. The first charge of $1.1 billion is for the resolution of tax disputes associated with our stores business in Italy, and the settlement of a lawsuit. This charge primarily impacts our international segment and is largely recorded in the fulfillment and other operating expense line items. The second charge is $730 million for the estimated severance costs.

This charge impacts all three of our segments and is recorded primarily in the fulfillment sales and marketing and technology and infrastructure expense line items. The third charge of $610 million is for asset impairments primarily related to physical stores. This charge primarily impacts the North America segment as reported in other operating expense line. Moving on to our segment results. In the North America segment, fourth quarter revenue was $127.1 billion, an increase of 10% year over year. International segment revenue was $50.7 billion, an increase of 11% year over year excluding the impact of foreign exchange. Worldwide paid units grew 12% year over year, because our highest quarterly growth rate in 2025.

The fourth quarter marked a strong finish to the year as we delivered for customers during the peak holiday season. Our sharp pricing, vast selection, and record vast delivery speeds resonated with customers. They appreciate the convenience of receiving their items quickly. Bring gifts for family and friends, to everyday essentials and perishable groceries. Our millions of global third-party sellers continue to be an important contributor to our broad selection. In Q4, worldwide third-party seller unit next was 61%. We continue to invest in tools and services including a comprehensive suite of AI tools that help our selling partners manage and grow their businesses.

Shifting to profitability, North America segment operating income was $11.5 billion with an operating margin of 9%. Up from an 8% margin in 2024. International segment operating income was $1 billion with an operating margin of 2.1%. Excluding the impact of special charges mentioned earlier, International segment operating margins also expanded year over year. We are pleased with the fulfillment network performance throughout the peak season. We made strong progress improving the cost structure of our network over the past few years. In The US, a regionalized network is operating at scale and we continue to make refinements. This regionalization has improved local inventory placement, leading to faster delivery and lower costs.

Last year, US Prime members received over 8 billion items the same or next day. Up more than 30% year over year. With groceries and everyday essentials making up half of the total items. For the third year in a row, globally in 2025, we achieved both our fastest ever delivery speeds for Prime members while also reducing our cost to serve. By leveraging our existing US network, we can now deliver perishable groceries to customers in more than 2,300 cities and towns, all with same day delivery. We saw significant adoption of this service throughout the year. When customers engage with our perishable offering, they demonstrate notably higher monthly spend compared to those who did not shop the category.

We also see the customer shopping perishable gross add three times more items to their same day delivery orders. Looking ahead, we see further opportunity to enhance productivity our global fulfillment network while delivering at faster speeds for customers. We'll continue optimizing inventory placement to drive down distance traveled, reduce touches per package, and improve package consolidation. As well as launch robotics and automation to increase efficiency and elevate the customer experience. Shifting to advertising. Advertising revenue grew 22% in the fourth quarter and we added over $12 billion of incremental revenue in 2025 alone. As our full funnel advertising approach of connecting brands with customers is resonating.

Simplifying the advertiser experience to enable brands to better reach customers wherever they are. Moving next to our AWS segment. Revenue was $35.6 billion, and growth accelerated to 24% year over year. We added $2.6 billion in quarter over quarter revenue and AWS now has an annualized revenue run rate of $142 billion. This acceleration was driven by both core and AI services as customers continue to modernize their infrastructure and migrate workloads to the cloud. Our AI offerings continue to resonate with customers. Including our argenic capabilities. This growth was helped in part by the more than one gigawatt of capacity we added in Q4.

In 2025, AWS added more data center capacity than any other company in the world. AWS operating income was $12.5 billion. We are seeing strong top line growth while remaining focused on driving efficiencies across the business. This includes investing in software and process improvements to optimize server capacity. Developing a more efficient network using our lower cost custom networking gear, advancing custom silicon. At the same time, we continue to rapidly develop products and services on behalf of customers. As we've long said, we expect AWS operating margins to fluctuate over time, driven in part by the level of investments we are making at any point in time. Turning to cash flows.

Our full year operating cash flow increased to $139.5 billion in 2025. Up 20% year over year due primarily to improved operating income and changes in working capital. Now turning to our Q1 financial guidance. Q1 net sales are expected to be between $173.5 billion and $178.5 billion. This guidance anticipates a favorable impact of approximately 180 basis points from foreign exchange rates. As a reminder, currencies can fluctuate during the quarter. Q1 operating income is expected to be between $616.5 billion and $21.5 billion. A few things to mention on the operating income guidance. Within the North America segment, we do expect a year over year cost increase of approximately $1 billion related to Amazon LEO.

We have more than 20 launches planned in 2026, and more than 30 in 2027 which means we are spending more on launching satellites each year. Select enterprise customers are testing Amazon LEO services now. And we expect a wider commercial rollout later this year. As a reminder, today, we do expense most of these LEO costs as incurred. We expect that later in the year, many of these costs, such as satellite manufacturing and launch services, will be capitalized. Within the international segment, we are continuing to invest more in our stores business to enhance the customer experience and to encourage retail demand to move online more quickly.

This includes bringing faster delivery options including Amazon Now or service which delivers to customers in thirty minutes or less. We are also working hard to stay sharp on pricing and seller fees. In our countries where we've had to be more aggressive to meet or beat competitors' prices. We like these investments because they will delay customers, grow our business, and we believe they will generate long-term positive return on invested capital. As we enter 2026, I'm energized by our team's strong execution. I want to thank everyone across the company for their hard work on behalf of our customers.

We remain focused on driving an even better customer experience which is the only reliable way to create lasting value for our shareholders. With that, let's move on to your questions.

Operator: At this time, we will now open the call up for questions. We ask each caller to please limit yourself to one question. If you would like to ask a question, please press star 1 on your telephone keypad. We ask that when you pose your question, you pick up your handset to provide optimum sound quality. Once again, to initiate a question, please press star then 1 on your touch tone telephone at this time. Please hold while we poll for questions.

Andrew Jassy: Thank you.

Operator: Our first question comes from the line of Mark Mahaney with Evercore ISI. Please proceed with your question.

Mark Mahaney: Okay. Thanks. I think, Brian, let me throw this to you or maybe to Andy. On the strong long-term return on invested capital, I think that's the debate in the market today. So could you give us a little bit more insight into how you think investors will be able to see that? Either, talk about the duration of the CapEx cycle that you're going through now or, what we should see in terms of profitability levels? And maybe also talk about, like, other de minimis or minimum free cash flow generation levels that you don't want to go below as you go through this CapEx cycle?

Just help us get to that get to your level of confidence in having a strong long-term return on that invested capital. Thank you.

Brian T. Olsavsky: Yeah. Sure, Mark. Thank you. I'll start from a financial side. So on the investments we are making, as Andy said earlier, you know, we are putting into service with customers all capacity that we are getting, and it's immediately useful. And we are also seeing a long arc of additional revenue that we see from other customers and backlog and commitments of people are anxious to make with us, especially for AI services. So you can see that it's working out its way into our P&L both through CapEx and also through our operating margin in AWS. AWS is 35% operating margin through Q4, up 40 basis points year over year.

Just talked about before that is going to fluctuate over time. It certainly has a headwind from the investments in AI and the depreciation on that CapEx. But we also work very hard to offset that with efficiencies. And cost reduction. So we will see how that develops over time. So but, yeah, we see long strong return on invested capital. We see strong demand for these services, and we continue to like the investments in this area. I would add to that. You know what? If you look at the capital, we are spending and intend to spend this year, it's predominantly in AWS.

And some of it is for our core workloads which are our non-AI workloads because they are growing at a faster rate than we anticipated. But most of it is in AI and we just have a lot of growth and a lot of demand. When you are growing 24% year over year with an annualized revenue run rate of $142 billion, you are growing a lot. And what we are continuing to see is as fast as we install this capacity, this AI capacity, we are monetizing it. So it's just a very unusual opportunity.

You know, as I've shared a lot of times I passionately believe that every customer experience that we know of today is going to be reinvented. With AI, there are gonna be a whole bunch of customer experience none of us ever imagined that are gonna become the norms of how we all operate every day and what we use. I think the other thing is that if you really want to use AI, in an expansive way, you need your data in the cloud and you need your application in the cloud. Those are all big tailwinds pushing people towards the cloud.

So we are gonna invest aggressively here, and we are gonna invest to be the leader in this space as we have been for the last number of years. We have, I think, a fair bit of experience over the years in AWS of forecasting demand signals and doing it in such a way that we don't have a lot of wasted capacity and that we also have enough capacity to serve the demand that's there. And I think we've also proven with AWS over the years how we build data centers, how we run them, and how we invent in there.

If you think about our chips and our hardware and networking gear and how we've invented and power that this isn't some sort of quixotic top line grab. You know, we have confidence that these investments will yield strong returns on invested capital. We've done that with our core AWS business. I think that will very much be true here as well and I think some of the things that you will see over time in the AI space is you're gonna keep seeing all the inference services, which is gonna be the majority of the long-term AI workloads is gonna be inference. You're gonna see the inference keep getting optimized. You're going to see higher utilization on those services.

You'll see prices normalize over a period of time. And then I think the companies that have not just the excellence and infrastructure but also the component that give them it'll give customers better price performance and give those companies themselves better economics are going to have advantaged financials. And I think if you look we're already off to a really good start having Traneum underneath some majority of our bedrock service and that's not just giving customers better prices but it also gives us better economics. So we see that following the same sorts of patterns we saw in the early days of our core AWS investment. I'm very confident we're gonna have strong return on invested capital here.

Operator: And the next question comes from the line of Douglas Anmuth with JPMorgan. Please proceed with your question.

Douglas Anmuth: Thanks so much for taking the questions. Can you just talk about how Project Ranir is running with Anthropic after first full quarter? And I think in the release, it talks about 500,000 ships, but a few months ago, you talked about getting to 1,000,000 as well. So if you could clarify that. And then maybe just a follow-up on Mark's question, are there any financial guardrails or governors in place that we should think about around the spend just in terms of operating income growth? Or positive free cash flow? Thanks.

Andrew Jassy: Yeah. I'll start with the training piece. We are very excited about the growth that we see in Traneum the future that we have there. I think if you look at what's happened in the early innings of AI over the first few years, you see a lot of usage but customers are really thirsty for better price performance. And Trainium has 30% to 40% better price performance than comparable GPUs, so it's very compelling to customers. You mentioned Project Rainier Anthropic is building their next they're training the next Claude model on top of Trainium two. And that's what Project Rainier is. We talked about 500,000 chips there. You will see that continuing to increase.

They're also using a fair bit of Trainium-two for other workloads and their own APIs beyond just Project Rainier but Trainium is a multibillion-dollar annualized run rate business at this point and it's fully subscribed. And what you're also seeing is Trainium three, which is the next version of Trainium, which we just started shipping. That's 40% more price performance than Trainium two. And we have there is very substantial amount of interest there. We expect that nearly all of that supply will be committed by somewhere around the middle of this year. And we're just in the process of building training four.

There's very substantial interest in training four which is coming in 2027 and we're already having conversations about training five so there is a lot of interest in training at this point and I think when you know, I think people know about our chips capability, our chips business, but I'm not sure folks realize how strong a chips company we've become over the last ten years. You know, if you look at what we've done with Trainium, if you look at what we've done with Graviton which is our CPU chip, which is about 40% better price performance than comparable x86 processors, 90% of the top 1,000 AWS customers are using Graviton very expansively.

If you combine Traneum and Graviton, it's well over a $10 billion annualized run rate business it's still very early there so I'm very optimistic about what we're seeing, know, the Project Rainier has gone very well. I think Anthropic is quite pleased with it. We've learned a lot in the process as well but it's early days with what's possible here. This is big business that's getting bigger and has a lot of potential. And I just briefly comment on your second question that you know, we are as I mentioned, this is what know, I think this is an extraordinarily unusual opportunity to forever change the size of AWS and Amazon.com, Inc. as a whole.

I think it also is an extraordinary opportunity for companies to change all their customer experiences and for startups to be able to build brand new experiences of businesses that would have taken much longer try to accomplish before that they can do right now. And so we see this as an unusual opportunity we are going to invest aggressively here to be the leaders because like we've been the last number of years and like I think we will be moving forward.

Operator: Thank you. The next question comes from the line of Ross Sandler with Barclays. Please proceed with your question.

Ross Sandler: Great. Andy, you mentioned if you called back how the AI market was currently a bit top heavy with a lot of the spend kinda clustering around a few of the AI native labs. So how is that changing as you look out into '26? And, specifically, how do you think you might extend your relationship with a company like OpenAI to maybe help Amazon.com's AI efforts both on the retail side and the AWS side. Thanks a lot.

Andrew Jassy: Yeah. The way I would describe what we see right now in the AI space is it's really kind of a barbelled market demand where on one end you have the AI labs who are spending gobs and gobs of compute right now along with and what I was consider a couple runaway applications. And then at the other side of the barbell, you've got a lot of enterprise who are getting value out of AI in doing productivity and cost avoidance types of workloads. These are things like customer service or business process automation or some of the fraud pieces. And then in that middle of the barbell are all the enterprise production workloads.

And I would say that the enterprises are in various stages at this point of evaluating how to move those, working on moving those, putting them into production. But I think that middle part of the barbell very well may end up being the largest and the most durable and I would put in the middle of that barbell too by the way I would put just the altogether brand new businesses and applications that companies build that right from the get go run-in production on top of AI.

And so I think that to me when I look at this and what's happening it's kind of unbelievable if you look at the demand of what you're seeing already with AI but the lion's share of that demand is still yet to come in the middle of that barbell and that will come over time, it will come as you have more and more companies with AI talent, as more and more people get educated with that AI background, as inference continues to get less expensive and that's a big piece what we're trying to do with Trainium and hardware strategy and as companies start to have success in moving those workloads to further and further success in moving those workloads run on top of AI.

So I think there's it's just a huge opportunity. It's still in the relative early stages even though it's growing at very unprecedented clip as we've talked about. You know, and then I think how do we see our relationships extending with other companies like OpenAI. You know, would tell you that this movement and what's happening in AI is it's very broad. It's gonna be a lot of companies. It is a lot of companies already. There's a number of AI labs but almost every company you talk to, almost every conversation we have on the AWS side starts with AI. And so we have very significant with a lot of different companies.

You know, I think we announced an agreement with OpenAI in November. We're excited about that agreement. It's a big one and we have a lot of respect for the company and we hope to continue to extend our partnership over time but this AI movement is not going to be a couple of companies. It's gonna be thousands of companies over time.

Operator: Thank you. The next question comes from the line of Michael Morton with MoffettNathanson. Please proceed with your question.

Michael Morton: Hi, good evening. Thank you for the question. This one's on the retail business. Andy, you've talked about how you're passionate. This is going to change experiences across the board. And you've shared some encouraging data points on Rufus. And we're seeing all the other Internet platforms roll out. Agentic protocols. I would love to see how you think this plays out for the retail business. And the on-site ads portion of the retail businesses. What seems like it could be a compression in the funnel as consumers get better answers over time. Anything there would be great. Thank you.

Andrew Jassy: I'm very optimistic about the customer experience that will ultimately be what customers use for agented shopping. And I think it's good for customers. I think it's going to make it easier for them. It's a big piece of why we've invested as significantly as we have in our own shopping assistant in Rufus. And if you haven't checked out Rufus recently, I really encourage you to do so. It's gotten much, much better and keeps getting better every month. And, you know, we have about three we have 300 million customers who use Rufus in 2025.

Customers who used Rufus are about 60% more likely to complete a purchase and so you just you know, you're seeing a lot of usage of it and a lot of growth and I think it's very useful. And, you know, I think at the same time, we will have relationships with third-party horizontal agents that can enable shopping as well. We have to collectively figure out a better customer experience. You know, it's still, you know, these horizontal agents don't have any of your shopping history, they get a lot of the product details wrong. They get a lot of the pricing wrong.

And so yeah, we have to try to find a customer experience together that better and, you know, a value exchange that makes sense for both parties. But I'm very hopeful that we'll get there over time.

We continue to have a number of conversations then I think you're gonna have to look at as time goes on you know, which types of you know, which shopping agents are consumers gonna use and it kinda reminds me in some ways of the early days of kinda all the search engines that were that were referring traffic to retailers and, you know, it's still a relatively small portion of the overall traffic and sales but of that fraction you have to ask how many consumers are going to prefer using a horizontal agent where it's kind of a middle person between the retailer and the consumer versus wanting to use a great agent from that retailer that has all its shopping history and that has all the data right there and makes it easy if you're just you know, spearfishing for something to shop for it right there or if you wanna do discovery, you can do it there and it's got the best data on shopping.

I think a lot of customers are ultimately gonna choose to use a great shopping agent from that retailer because if you think about what consumers really want in retail and retailer they want really broad selection, They want low prices. They want really fast delivery. And then they want they want a retailer that they can trust and that takes care of them. And I think horizontal agents are pretty good at aggregating selection but retailers are much better at doing all four of those items.

And so I'm very optimistic that people will use our shopping agent it's off to a great start I also expect that we'll work with other third-party agents over time as we work on the issues I mentioned earlier.

Operator: Thank you. The next question comes from the line of Brian Thomas Nowak with Morgan Stanley. Please proceed with your question.

Brian Thomas Nowak: Thanks for taking my question. Andy, I want to ask you one about the global retail business. This year. I know there's a lot of areas of investment in it that you're talking about. To sort of make improve the service, make it more durable over the long term, etcetera. But I'm assuming there are also sources of efficiency you expect to see this year. So can you sort of help us understand both sides of the ledger on retail this year? Where are some of the areas where you see the potential for sources of efficiency and cost to serve savings?

And then where should we be thinking about the areas of investment to drive more durable growth, robotics, etcetera? How does that sort of break down?

Andrew Jassy: Yeah. So I would say on the side of continuing to invest, to keep growing the retail business, you know, the kind of core drivers of demand continue to be the same. You know, we are gonna work really hard to expand selection and you've seen what we've done over the last several years, you know, the expansion of selection has been broad. And you'll see it on both ends of the spectrum. You know, we have a lot more of those luxury brands that have built presences in Amazon.com, Inc. had success found that we could we could manage their brand presentation the right way and they've been very happy.

Mean you only have to look at L'Oreal as an example too of just how fast that business is growing and how happy our partners have been and at the same time we are working really hard to continue to expand the amount of everyday essentials that we offer our customers. And the growth in everyday essentials in our business is really remarkable as I mentioned in my opening comments.

And one out of three units now that we that we move are everyday essentials and what we find there is that the more the customers can rely on us for everyday essentials and the lower ASP items, they just choose to do more their downstream shopping with us in every way. We're just more front of mind. And so, yeah, I think a big piece of why we have captured more and more of those everyday essentials and you see it also in our grocery business with perishables too.

It's just our speed of delivery improvements over the last three years has been really market mean, it's customers it's the one thing I get stopped on the street most often about which is I just can't believe how quickly from when I order something I get it my door and how reliable you are. I think along that speed of delivery piece it's also quite interesting what's happening with QuickCommerce, you know, and we have this offering called Amazon that we've largely started outside The US and India and The UAE and that gets thousands of items to customers within thirty minutes and it really is it's quite interesting how quickly that is growing.

And I think that it's just another one of those things like everyday essentials that when you're when you're able to order more and more from Amazon.com, Inc., you just think of Amazon.com, Inc. first if it's a great experience that we're offering for whatever you're looking for. But in our if you look in India, which is the place we've rolled out QuickCommerce the fastest, customers who try QuickCommerce are shopping with triple the frequency than they did before they tried us in Quick. So those are all areas I think are pretty exciting that we're expanding.

You'll see us continue to expand what we're doing in the perishable side too which we're quite excited about and we are able to deliver perishables same day in thousands of cities around the world now. In the cities in which we have those perishables available, nine out of 10 top items that are ordered in that geography are perishables. So we're just having a lot of success with that too and people buy perishables from us. After they buy perishables, they're shopping with us twice as frequently. So lot of good things to like there. And then, you know, on the efficiencies, we are I mean, we always have a very long list of these.

That we're working on, Brian. And, you know, it's true today as well. If you look even you know, I mentioned I talked a lot about regionalization. In our fulfillment network, particularly in The US. You know, over the last couple of years. And I said we weren't done honing that and that's true. It's just we, you know, we don't talk about it every time but if you look at what we've done there we've extended the number of regions. You know, it was eight. It's now 10. We've extended regionalization to what we do with our inbound delivery to be much more efficient, being able to get more items closer to customers more quickly.

We have made a lot of we're doing a lot of work and we've made a huge amount of progress in being able to get more units into each box. And as we're able to get more units into each box it obviously saves shipments and we drive better operating income when we do that and we've made very significant progress there but have a lot more planned. It's part of by the way that

improvement is part of what helps us do things like I was talking about earlier in adding to a delivery in near real time. And then, you know, robotics, as you mentioned, is another big one for us.

You know, we have over a million robots today in our fulfillment network. They take care of all sorts of functions, but still a fraction of what I think we are going to be able to enable over time which will allow our you know, we'll always have a lot of people that we employ in our fulfillment network, but they'll you know, they'll leave to the robotics things that, you know, that are more repetitive so it's better productivity for the business, more safe for teammates and there's real cost efficiencies in that as well. So, a lot on both sides of the ledger as always.

Operator: Thank you. And our final question comes from the line of Eric Sheridan with Goldman Sachs. Please proceed with your question.

Eric Sheridan: Thanks so much for taking the question. Maybe a few parts just on AWS. Can you speak to the current state of your revenue backlog as of Q4? And also discuss a little bit about what you see both for internal use cases and external client needs with respect to any imbalance between supply and demand around AI efforts? And how you think about closing the gap on those as more capacity comes online through 2026? Thank you.

Andrew Jassy: Yeah. That's a lot of parts. I'll start with the first one. Which is on backlog, our backlog is $244 billion. That's up 40% year over year. I think it's up 22% quarter over quarter. You know, we have and we have a lot of deals that are in the pipeline. There's just a as I mentioned earlier, there is a lot of demand for AWS right now. In the AI space and also in the core AWS space. Your second question was internal and external use cases. I would and then the impact around supply and demand. You know, the vast majority of our the capital that we spend and the capacity that we have consumed by external customers.

We have all Amazon.com, Inc. has always been a very large AWS customer, very helpful AWS customer because they're very demanding and they use the services very expansively and stretch the limits as we launch things. So they've always been a very important big customer but always a very small fraction of the total and that's true today in AI as well as the overall AWS business. You know, on internally, we have all sorts of ways that we are using AI.

We have over a thousand AI applications that we've either deployed or in the process of building, and they range from our shopping assistant in Rufus that we were just talking about to Alexa Plus, which is a really large scale generative AI application. To applications in our fulfillment network that allow us to have more forecasting predictions to how we do customer in our customer service chatbot to how we are making it much easier for brands to create advertisements and to optimize all their campaigns or across the full funnel of advertising options we have.

To, you know, in live sports, you watch Thursday night football, can see you know, defense of alerts which predict which player is gonna blitz or, you know, pocket health. I mean, we in every one of our businesses, you see a very broad use of AI to improve the customer experience. And in many cases, just to completely reinvent was possible before. I mean, it's pretty neat to use something like Lens where you may see something you want to buy, you can just take a picture of it in the app and it finds the item on the detail page and buy it one click. Kinda magic.

You know, externally, I would say it's kind of what I said earlier. You have AI labs, you know, consuming lots and lots of capacity. Both for training as well as for the inference and the research across what they're doing with their different applications and models. We see enterprises, all sorts of workloads, you know, customer service automation, business process automation, fraud, completely reinventing their application you know, agentic coding applications, legal applications, Suno is a really cool example. An AWS customer that's kind of reinvented how you can write music and build music.

So really across the board and, you know, and I just think on the supply and demand, what I would tell you is you know, we're growing 24% year over year a $142 billion annualized run rate business. So we're growing at really an unprecedented rate. Yeah. I think every provider would tell you, including us, that we could actually grow faster if we had all the supply that we could take. And so we are being incredibly scrappy around that. If you look know, in the last twelve months, we added 3.99 gigawatts of power. Just for perspective, that's twice what we had in 2022 when we were in $80 billion annual run rate business.

We expect to double it again by the '27. We added 1.2 gigawatts of power in Q4 just quarter over quarter. So it's so we are our team is being aggressive scrappy and inventive and adding capacity as fast as we can. I you know, we'll add a lot more in '26 and '27 and in '28 for that matter. And we're very optimistic we can continue to grow in the ballpark of what we have.

Dave Fildes: Thanks for joining us on the call today and for your questions. A replay will be available on our Investor Relations website for at least three months. We appreciate your interest in Amazon.com, Inc., and we look forward to talking with you again next quarter.